



PASS 2023 Pre-conference Workshop

Azure SQL and SQL Server: All Things Developers



PASS 2023 Pre-conference Workshop

Azure SQL and SQL Server: All Things Developers





Today's Speakers



Anagha Todabagi
Microsoft



Brian Spendolini
Microsoft



Schedule

Pre-Conference Sessions

8:30 AM – 10:30 AM

Refreshment Break

10:30 AM – 10:45 AM

Pre-Conference Sessions

10:45 AM – 12:15 PM

Pre-Conference Lunch

12:15 PM – 1:15 PM

Pre-Conference Sessions

1:15 PM – 2:45 PM

Refreshment Break

2:45 PM – 3:00 PM

Pre-Conference Sessions

3:00 PM – 4:30 PM

Agenda

- Workshop Overview
- Workshop Tasks
- Various Breaks and Lunch

Workshop Goals

- Get comfortable with new development tools
 - Local and cloud
 - New CLIs
- Expose you to REST/GraphQL
- Get started with Azure SQL Database
- Developing with Azure
- Expand your thinking about app development
- No marketing – all hands on

WORKSHOP OVERVIEW



Workshop Overview/Tasks

- Setup your environment
- Create database and objects
- Use Data API builder
- Applications with SWA
- Deploy to Azure SQL Database
- Invoke REST endpoints in the database
- Create a change data stream
- CICD with GitHub Actions

Getting Started

Accounts Needed

- GitHub account
- Azure account
 - Free Azure SQL Database

Getting Started

Development Environment



GitHub



docker



Static
Web Apps
&
Data API Builder



sqlcmd



Azure
Functions

Getting Started

Development Environment



GitHub



docker



SQL



sqlcmd



Azure
Functions



Static
Web Apps
&
Data API Builder



Azure
Functions



OpenAI



CHAPTER 1: SETUPS



Chapter 1: Setups

- Fork the repository
- Create the codespace

<https://aka.ms/forkit>



Code

Issues

Pull requests

Actions

Projects

Wiki

Security 2

Insights

Settings



azure-sql-db-developers-workshop

Public

Edit Pins

Watch 72

Fork 20

Star 3

main

1 branch

3 tags

Go to file

Add file

Code

Brian Spendolini and Brian Spendolini v1.5

✓ 07adcfc yesterday 196 commits

.devcontainer	v1.5	2 weeks ago
.github	.github/PULL_REQUEST_TEMPLATE.md committed	5 months ago
app	v1.5	3 weeks ago
client	v1.5	3 weeks ago
docs	v1.5	yesterday
labFiles	v1.5	2 weeks ago
scripts	v1.5	last week
CHANGELOG.md	v1.5	yesterday
CONTRIBUTING.md	v1.5	last week
LICENSE.md	LICENSE.md committed	5 months ago
README.md	v1.5	2 weeks ago
swa-cli.config.json	v1.5	3 weeks ago

About

A workshop for developing with the Azure SQL Database and Azure Services

Readme

MIT license

Code of conduct

Activity

3 stars

72 watching

20 forks

Report repository

Releases 3

v1.5 Latest

6 minutes ago

+ 2 releases

Packages

Setups

Repository Contents

.devcontainer

- Codespace docker config

client

- contains the Todo application for use with Data API builder and SWA

app

- contains the sample JavaScript test application for use with Data API builder and SWA

docs

- Has all the chapters/lessons of the workshop

scripts

- Installs all the needed dev tools

labFiles

- contains files to be used with the workshop in various chapters

GitHub Codespaces

- A development environment that's hosted in the cloud
- Each codespace is hosted by GitHub in a Docker container
- Browser based environment
- Config files are contained in the repository
- VS Code extension available

This branch is up to date with Azure-Samples/azure-sql-db-developers-workshop:main.

Contribute ▾ Sync fork ▾

JetterMcTedder v1.5

eadcad5 yesterday 167 commits

.devcontainer	v1.5	last week
.github	.github/PULL_REQUEST_TEMPLATE.md committed	4 months ago
app	v1.5	5 days ago
client	v1.5	5 days ago
docs	v1.5	yesterday
labFiles	v1.5	5 days ago
scripts	v1.5	2 days ago
CHANGELOG.md	v1.5	2 days ago
CONTRIBUTING.md	v1.0	4 months ago
LICENSE.md	LICENSE.md committed	4 months ago
README.md	Updated README	3 weeks ago
swa-cli.config.json	v1.5	last week

README.md



About

A workshop for developing with the
Azure SQL Database and Azure Services

Readme

MIT license

Activity

0 stars

0 watching

21 forks

Releases

No releases published

[Create a new release](#)

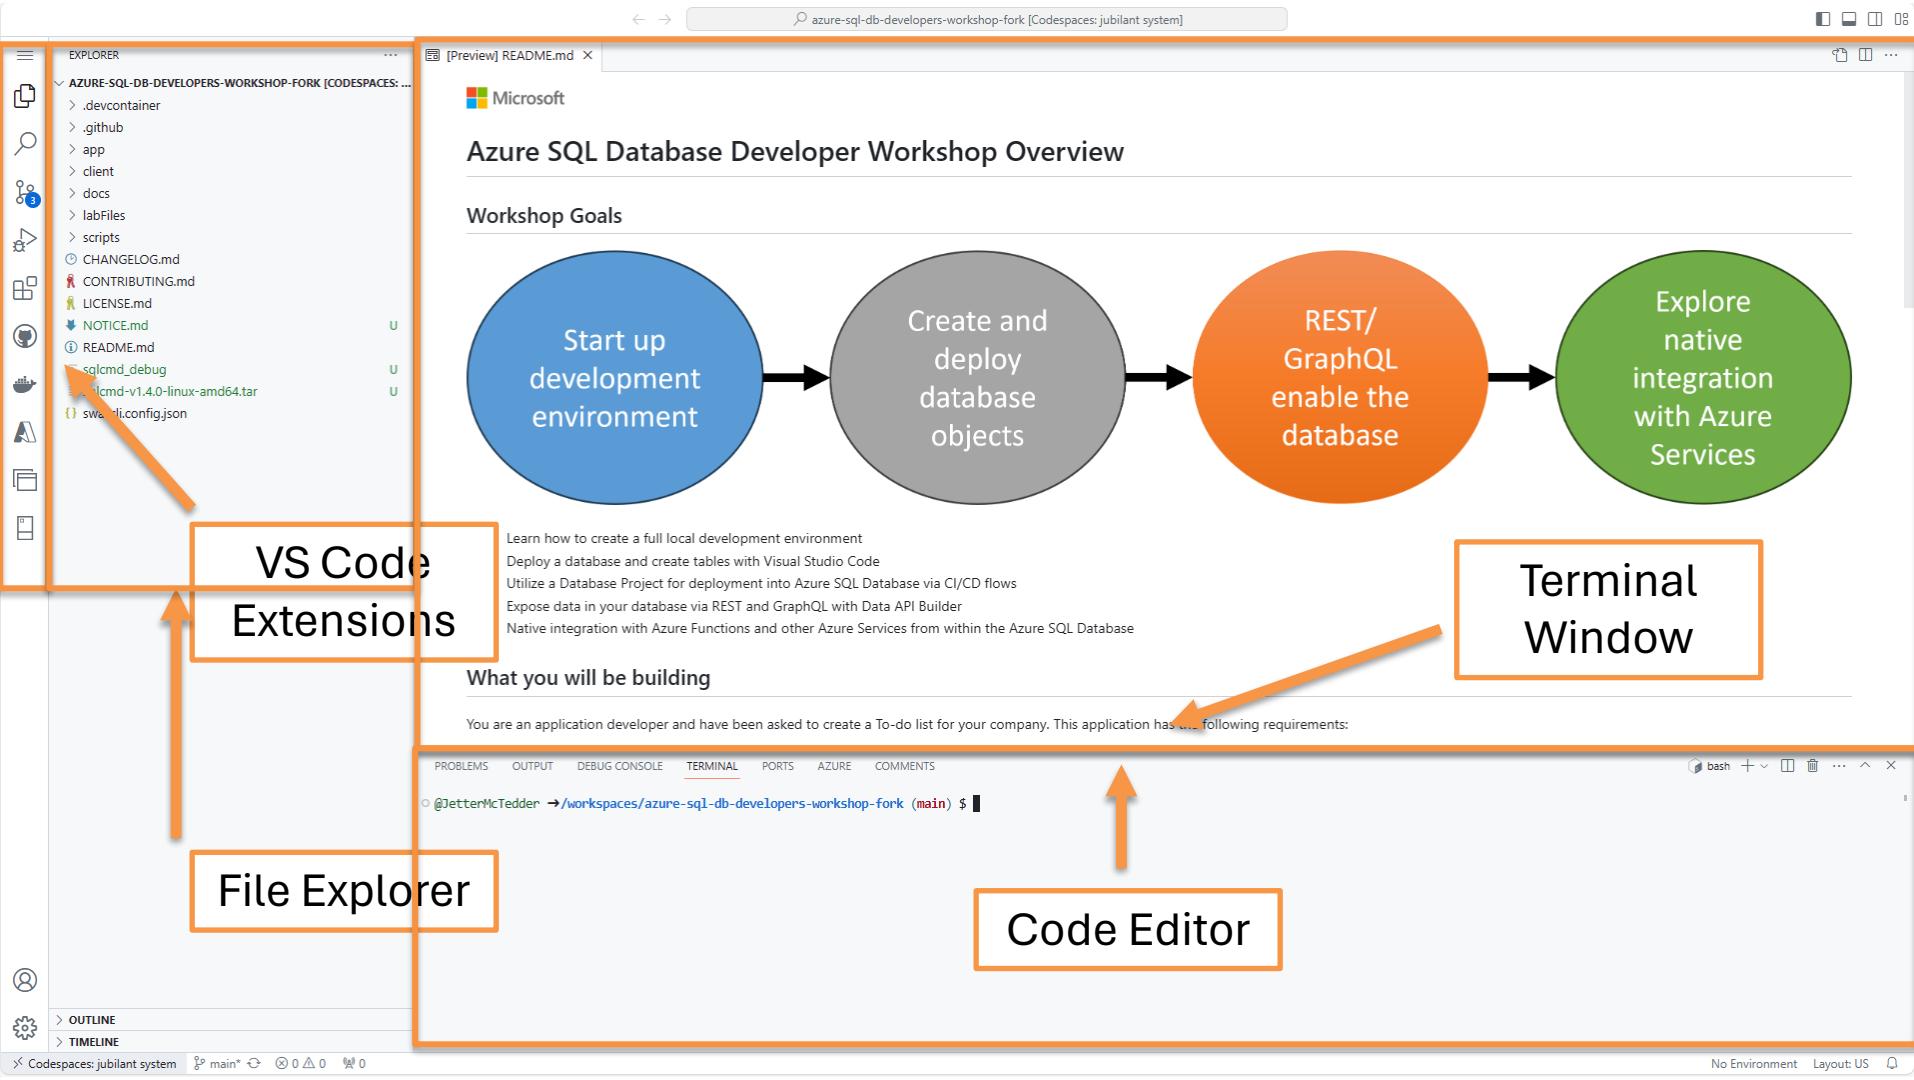
Packages

No packages published

[Publish your first package](#)

Languages





This screenshot shows a Microsoft Visual Studio Code (VS Code) interface for an Azure SQL DB Developers Workshop workspace. The workspace is titled "azure-sql-db-developers-workshop [Codespaces: super space disco]".

Explorer View: On the left, the Explorer view shows the project structure under "AZURE-SQL-DB-DEVELOPERS-WORKSHOP...". It includes a ".devcontainer" folder, a ".github" folder, a ".vscode" folder with 48 files, an "app" folder, a "client" folder, and a "database/devDB" folder containing "bin", "obj", and several SQL files: "address.sql", "delete_todo.sql", "devDB.sqlproj", "get_person_by_pet.sql" (selected), "insert_todo.sql", "person.sql", "README.md", "Script.PostDeployment1.sql", "todo.sql", and "update_todo.sql". There are also "docs", "labFiles", "scripts", "swa-db-connections", "triggerBinding", "CHANGELOG.md", "CONTRIBUTING.md", and configuration files like "dab-config.json", "OUTLINE", and "TIMELINE".

Terminal View: The bottom right terminal window shows the output of the command "dab --version", which displays the version of Microsoft.DataApiBuilder as 0.8.52+c69925060e1942d28515b9c4b89ea24832da0c7c.

Code Editor: The main editor area displays the content of the "get_person_by_pet.sql" file, which contains a CREATE PROCEDURE definition:

```
1 CREATE PROCEDURE dbo.get_person_by_pet
2     @pet nvarchar(100)
3 AS
4 BEGIN
5     select *
6     from dbo.person
7     where pet_preference = iif(NULLIF(@pet, '') IS NOT NULL,@pet,pet_preference);
8 END;
9 GO
```

Utilities Check

Checks to run once the post-build scripts are finished:

Functions Core Tools

- `func --version`

Data API builder

- `dab --version`

sqlcmd

- `sqlcmd --version`

Static Web Apps

- `swa --version`

```
@JetterMcTedder → /workspaces/azure-sql-db-developers-workshop (main) $ func --version
4.0.5390
@JetterMcTedder → /workspaces/azure-sql-db-developers-workshop (main) $ dab --version
Microsoft.DataApiBuilder 0.8.52+c69925060e1942d28515b9c4b89ea24832da0c7c
@JetterMcTedder → /workspaces/azure-sql-db-developers-workshop (main) $ sqlcmd --version
v1.4.0
Legal docs and information: aka.ms/SqlcmdLegal
Third party notices: aka.ms/SqlcmdNotices
@JetterMcTedder → /workspaces/azure-sql-db-developers-workshop (main) $ swa --version
1.1.4
@JetterMcTedder → /workspaces/azure-sql-db-developers-workshop (main) $ □
```

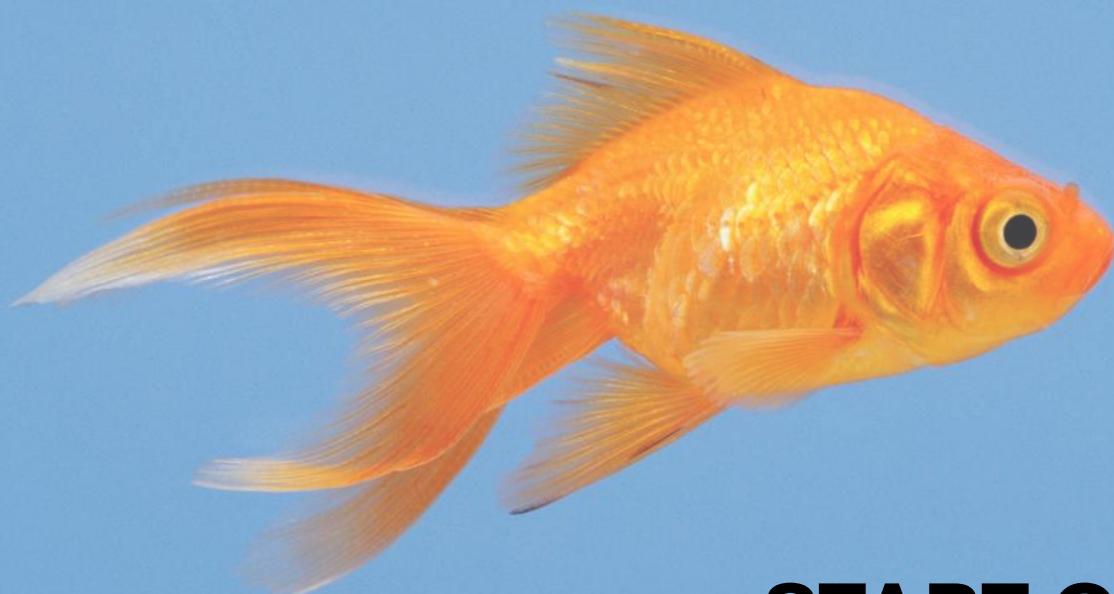
Knowledge Check

What cloud-based development tool are you going to use today?

- A. GitHub CICD
- B. Azure DevOps
- C. GitHub Codespaces
- D. Microsoft Teams
- E. A and D

Chapter 1: Setups

- Fork the repository
- Create the codespace



START CHAPTER 1



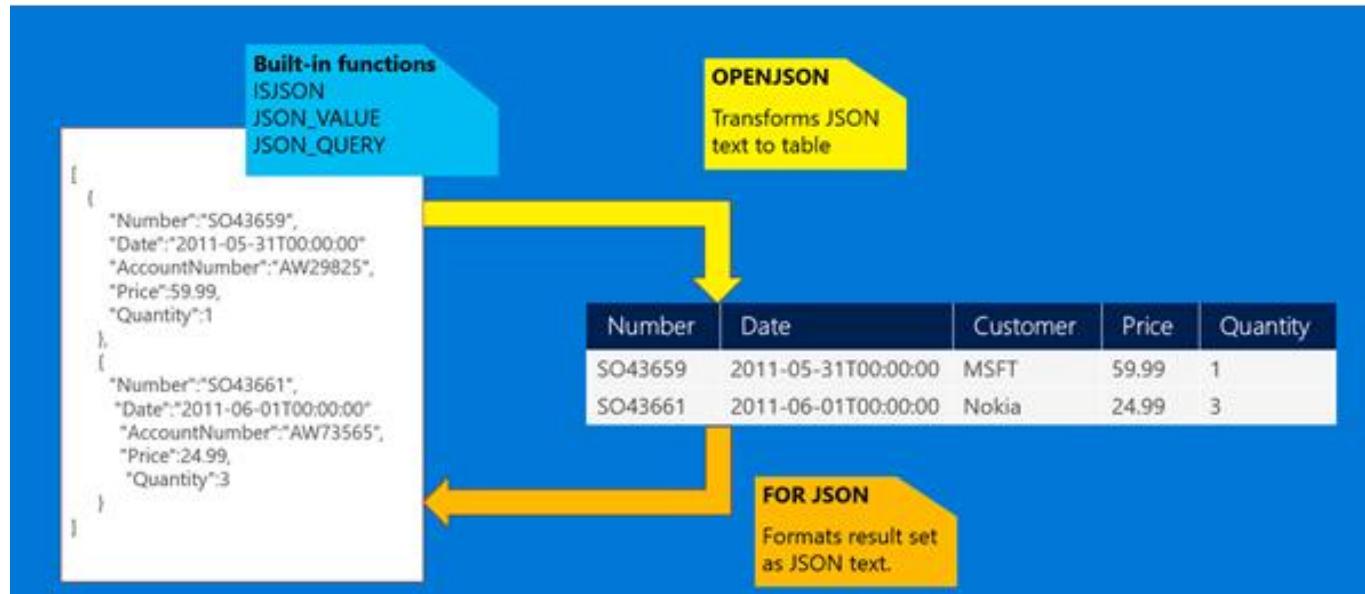
CHAPTER 2: CREATE A DATABASE AND OBJECTS

Azure SQL Database

- Features for developers
 - Multi-model support
 - Relational / JSON / Graph / Geospatial / XML
 - Row-Store and Column-Store
 - on the same table
 - Ledger Tables
 - Replicas / High-Availability / Scale-Out / Scale-Up

JSON Support

Azure SQL fully support storage, manipulation and creation of JSON documents



<https://docs.microsoft.com/sql/relational-databases/json/json-data-sql-server>

JSON Support

ISJSON: Tests whether a string contains valid JSON.

JSON_PATH_EXISTS: Tests whether a specified SQL/JSON path exists in the input JSON string.

JSON MODIFY: Updates the value of a property in a JSON string and returns the updated JSON string.

JSON_VALUE: Extracts a scalar value from a JSON string.

JSON Support

JSON_QUERY: Extracts an object or an array from a JSON string.

JSON_OBJECT: Constructs JSON object text from zero or more expressions.

JSON_ARRAY: Constructs JSON array text from zero or more expressions.

FOR JSON: Converts SQL Server data types to JSON types.

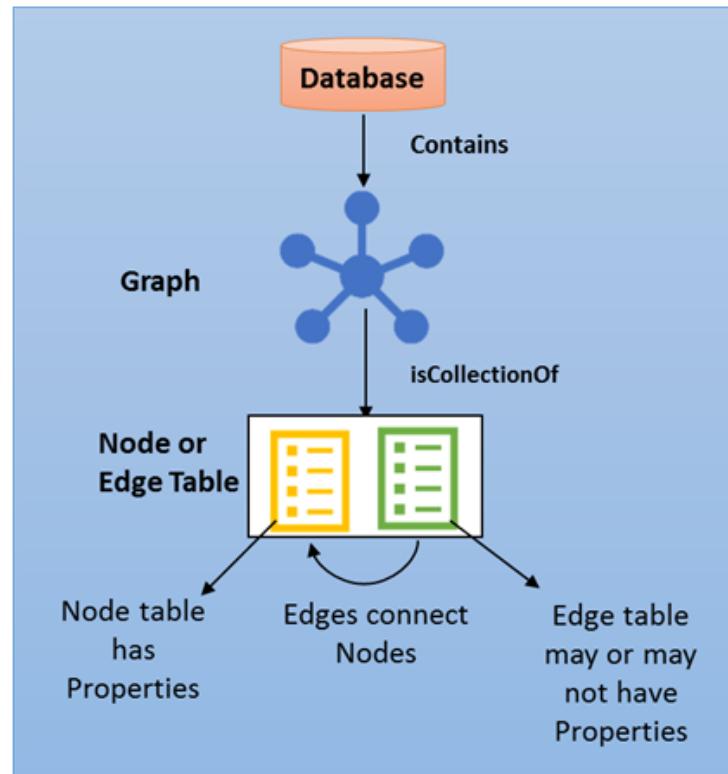
OPENJSON: Converts JSON text into a set of rows and columns.

JSON Support

```
SELECT Name,  
       Surname,  
       JSON_VALUE(jsonCol, '$.info.address.PostCode') AS PostCode,  
       JSON_VALUE(jsonCol, '$.info.address."Address Line 1"')  
         + ' ' + JSON_VALUE(jsonCol, '$.info.address."Address Line 2"')  
           AS Address,  
       JSON_QUERY(jsonCol, '$.info.skills') AS Skills  
FROM People  
WHERE ISJSON(jsonCol) > 0  
      AND JSON_VALUE(jsonCol, '$.info.address.Town') = 'Belgrade'  
      AND STATUS = 'Active'  
ORDER BY JSON_VALUE(jsonCol, '$.info.address.PostCode');
```

Graph Models

- Create EDGE and NODE objects
 - For simplicity they look like tables
- Use MATCH to find rows that satisfy some relationship condition
 - Syntax is close to Cypher language (supported by Neo4J)
- Getting started:
 - <https://aka.ms/azuresqlgraph>



Graph Models

```
CREATE TABLE Person (
    ID INTEGER PRIMARY KEY,
    Name VARCHAR(100),
    Age INT)
AS NODE;
```

```
CREATE TABLE Friends (
    StartDate date)
AS EDGE;
```

--Find friends of John

```
SELECT Person2.Name
FROM Person Person1,
     Friends,
     Person Person2
WHERE MATCH(Person1-(Friends)-
>Person2)
AND Person1.Name = 'John';
```

Graph Models

The screenshot shows the Neo4j Studio interface with the following components:

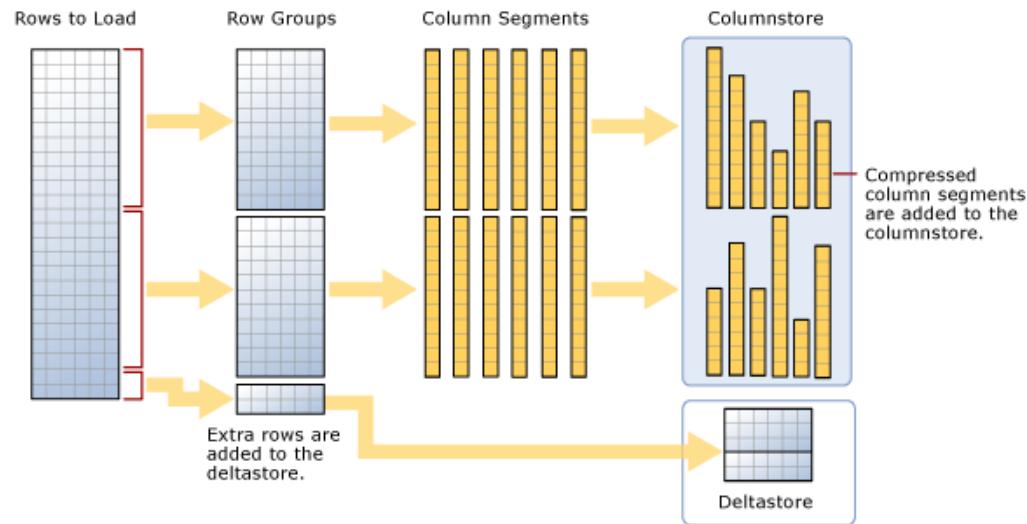
- Node Properties:** A section above the query editor containing two arrows pointing down towards the code area.
- Nodes that this edge connects:** A section above the results table containing three arrows pointing down towards the results table.
- Edge Properties:** A section above the results table containing one arrow pointing down towards the results table.
- Code Area:** Displays a Cypher query:

```
1 SELECT Person2.Name Untitled-2
2
3   SELECT Person2.Name
4   FROM Person Person1,
5       friendOf,
6       Person Person2
7   WHERE MATCH(Person1-(friendOf)->Person2)
8   AND Person1.Name = 'John';
```
- Results Area:** A table titled "RESULTS" showing the output of the query:

	Name
1	Mary
- Tables Labels:** Below the interface, there are labels for the tables:
 - Person Node Table
 - Friends Edge Table

Columnstore

- Store data by columns instead of by rows
- Great for analytical queries
 - Aggregations / Projections
- Columnstore is updatable
 - available from 2016
 - “Deltastore”



Reasons for Columnstore Indexes

- Columns store values that commonly have similar values, which result in high compression rates.
- High compression rates improve query performance by using a smaller in-memory footprint
- Queries often select only a few columns from a table, which reduces total I/O from the physical media.

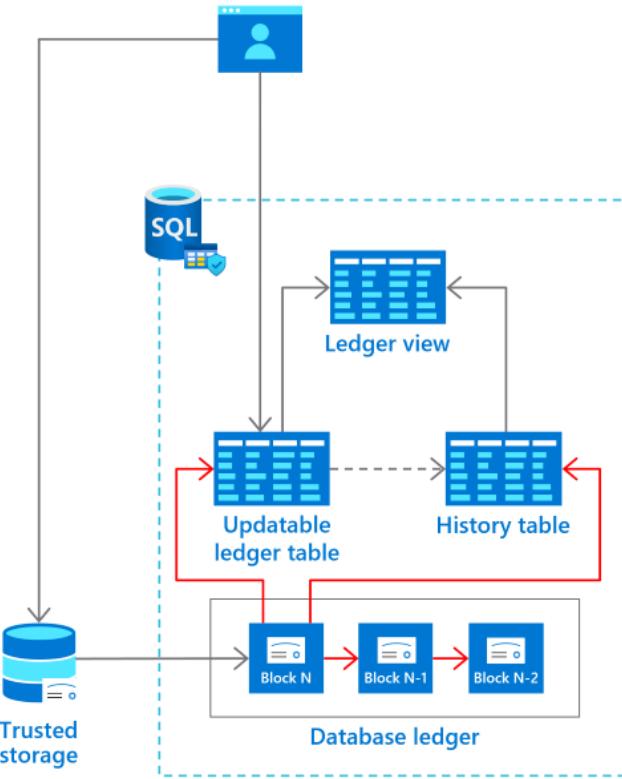
Ledger Tables

Updatable ledger tables

Updatable ledger tables are system-versioned tables on which users can perform updates and deletes while also providing tamper-evidence capabilities.

When updates or deletes occur, all earlier versions of a row are preserved in a secondary table, while the latest version of the row remains in the ledger table.

Ideal for systems/tables where auditing and record accuracy are vital.



Ledger Tables

```
CREATE TABLE [dbo].[HighScores]
(
    [ScoreID] INT IDENTITY,
    [PlayerID] INT NOT NULL,
    [FirstName] VARCHAR (50) NOT NULL,
```

CommitTime	UserN...	Sco...	Pla...	Firs...	LastNa...	Game	Score	Date	Operation
2023-04-20 23:20:49...	bill	2	2	Bill	Mickelson	Horse Monkey	11099913	2023-04-20	INSERT
2023-04-20 23:20:49...	bill	2	2	Bill	Mickelson	Horse Monkey	11099911	2023-04-20	DELETE
2023-04-20 23:20:35...	sqladmin	1	1	Steve	Wobble	Horse Monkey	11099912	2023-04-20	INSERT
2023-04-20 23:20:35...	sqladmin	3	3	John	Hill	Universe Invaders	218870	2023-04-20	INSERT
2023-04-20 23:20:35...	sqladmin	2	2	Bill	Mickelson	Horse Monkey	11099911	2023-04-20	INSERT

LEDGER = ON);

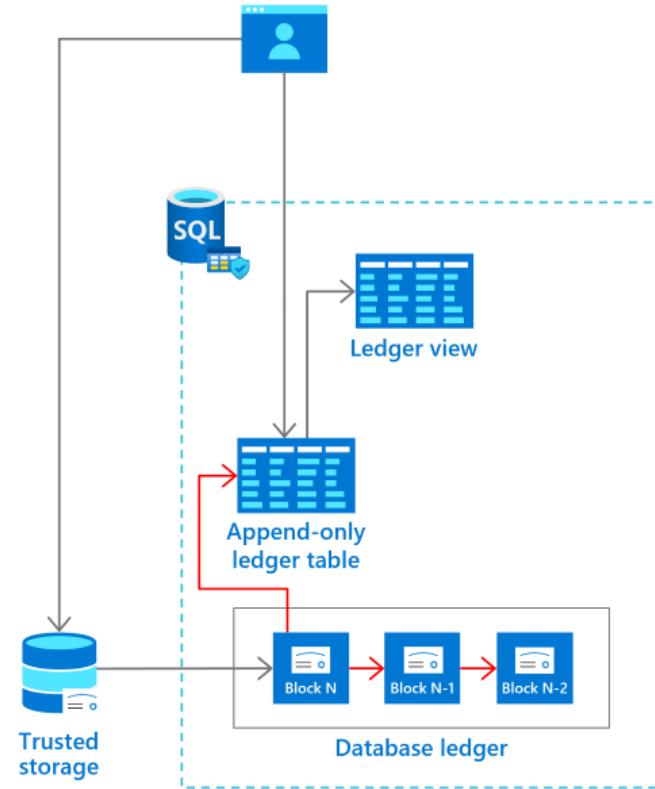
```
INSERT INTO [dbo].[HighScores]
VALUES      (1, 'Steve', 'Wobble', 'Horse Monkey', 11099912, GETDATE()),
            (2, 'Bill', 'Mickelson', 'Horse Monkey', 11099911, GETDATE()),
            (3, 'John', 'Hill', 'Universe Invaders', 218870, GETDATE());
```

Ledger Tables

Append-only ledger tables

Append-only ledger tables allow only **INSERT** operations on your tables, which ensure that privileged users such as database administrators can't alter data through traditional Data Manipulation Language operations.

Ideal for security information event and management systems such as badge in/out or ID readers.



Ledger Tables

```
UPDATE [dbo].[ArcadeRoomAccess]
```

```
SET Timestamp = GETDATE()
```

```
WHERE AccessID = 4;
```

Msg 37359, Level 16, State 1, Line 1

Updates are not allowed for the
append only Ledger table 'dbo.ArcadeRoomAccess'.

Azure SQL DB Hyperscale

Azure SQL Database Hyperscale provides the flexibility and performance modern applications require:

- Independently scalable compute and storage, without sacrificing performance
- Rapid, auto-scaling storage support up to 100TB
- Low latency, high throughput
- Snapshot-based backups that do not impact query performance
- Available serverless compute and resource-sharing elastic pools

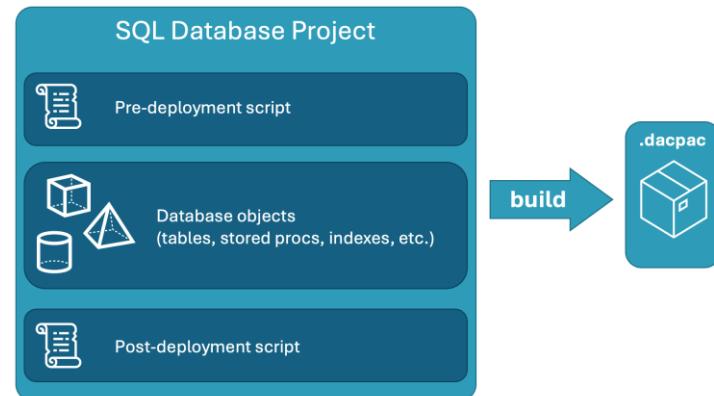
Now with lower pricing – save up to 35%

Chapter 2: Database Tasks

- Create a database SQL project
- Create a SQL 2022 database in the codespace
- Create database objects
- Publish objects to the local SQL 2022 database

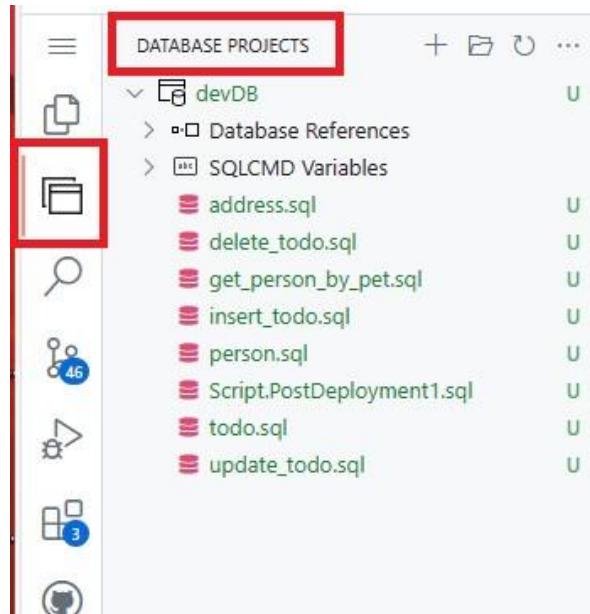
SQL Database Projects

- An Azure Data Studio and Visual Studio Code extension
- For developing SQL databases and objects in a project-based development environment
- Update/Create project from database
- Compare schemas
- Publish all or just changes
- Participate in CI/CD flows
- Deployment time configurations



SQL Database Projects

- `dotnet new sqlproj -n "devDB" -tp "SqlAzureV12"`



Create a SQL Database

- SQLCMD CLI
 - Now written in Go
- Use the go-mssql driver
- More Azure Active Directory (Entra ID) authentication options
- New commands
 - `sqlcmd create mssql --using https://url.com/backup.bak`

Create a SQL Database

- `sqlcmd create mssql -u devDB --accept-eula`
- `sqlcmd config connection-strings`

```
● @JetterMcTedder → /workspaces/azure-sql-db-developers-workshop/database (main) $ sqlcmd create mssql -u devDB --accept-eula
Downloading mcr.microsoft.com/mssql/server:latest
Starting mcr.microsoft.com/mssql/server:latest
Created context "mssql" in "/home/vscode/.sqlcmd/sqlconfig", configuring user account...
Disabled "sa" account (and rotated "sa" password). Creating user "vscode"
Creating default database [devDB]
Now ready for client connections on port 1433
```

HINT:

1. Run a query: `sqlcmd query "SELECT @@version"`
2. Start interactive session: `sqlcmd query`
3. View sqlcmd configuration: `sqlcmd config view`
4. See connection strings: `sqlcmd config connection-strings`
5. Remove: `sqlcmd delete`

Create Database Objects

- Tables
 - ID columns
 - Foreign/Primary Keys
- Table APIs - Stored procedures for
 - Insert
 - Update
 - Delete

Person

Name	Data Type
person_id	int (IDENTITY)
person_name	nvarchar(200)
person_email	nvarchar(200)
pet_preference	nvarchar(100)

Address

Name	Data Type
address_id	int (IDENTITY)
person_id	int
address	nvarchar(200)

FK

Todo

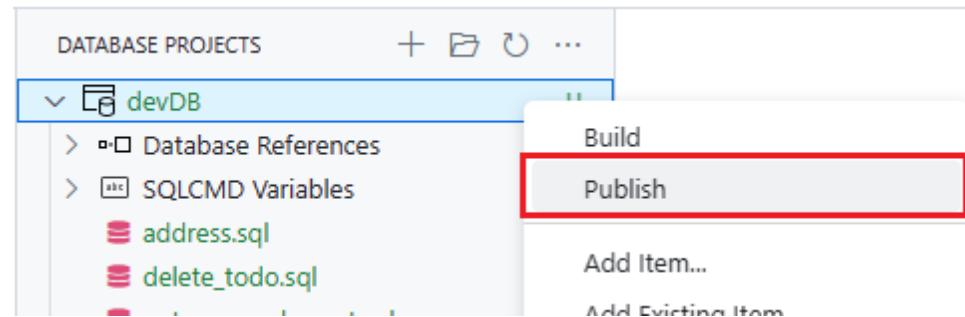
Name	Data Type
id	uniqueidentifier
title	nvarchar(1000)
completed	bit
owner_id	nvarchar(126)
position	int

Products

Name	Data Type
product_id	int
product_name	nvarchar(1000)
ListPrice	money

Publish objects to the local database

- Use SQL Database Projects
 - Publish to the local database created with sqlcmd
- Deploy the tables, stored procedures and run a post deploy script
- Verify the deployment



Testing the Table APIs Reminder

You can also test out the stored procedures with the following code:

```
exec dbo.insert_todo @title = 'My Test Todo', @owner_id = '1001001', @order = 1;
select * from dbo.todo;
GO
```

Copy and paste the resulting ID from the result set into the area marked **COPIED ID FROM RESULT SET** and use it in the next stored procedure example:

```
exec dbo.update_todo @id = 'COPIED ID FROM RESULT SET', @title = 'zzzz', @owner_id =
'1001001';
select * from dbo.todo;
GO
```

And again, copy and paste the resulting ID from the result set into the area marked **COPIED ID FROM RESULT SET** and use it in the next stored procedure example:

```
exec dbo.delete_todo @id = 'COPIED ID FROM RESULT SET', @owner_id = '1001001';
select * from dbo.todo;
GO
```

Knowledge Check

Which of the following can be used in Azure SQL Database

- A. JSON
- B. XML
- C. Geospatial
- D. Graph
- E. All of the above

I have an application that logs cars entering and exiting a secured area, which ledger table should I use?

- A. Updatable
- B. Insert Only
- C. Append Only
- D. Delete Updatable
- E. A, B, and D

Knowledge Check

Are SQL Database Projects good for CI/CD?

A. Yes

B. No

Chapter 2: Database Tasks

- Create a database SQL project
- Create a SQL 2022 database in the codespace
- Create database objects
- Publish objects to the local SQL 2022 database

START CHAPTER 2



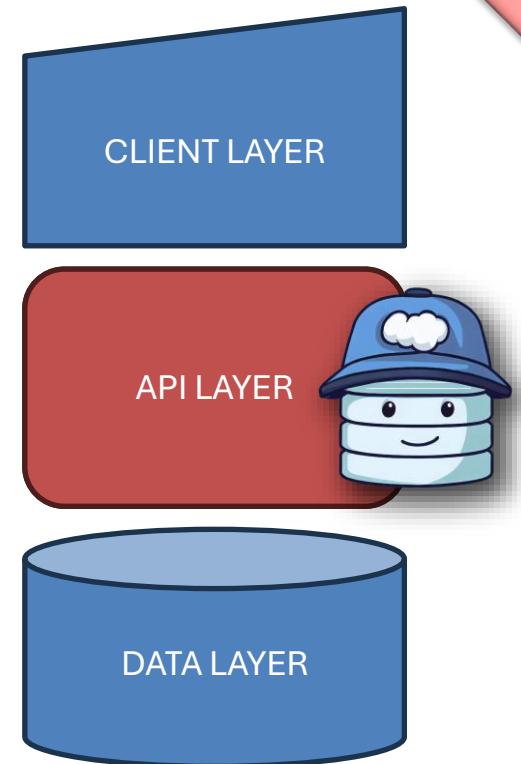
CHAPTER 3:

DATA API BUILDER



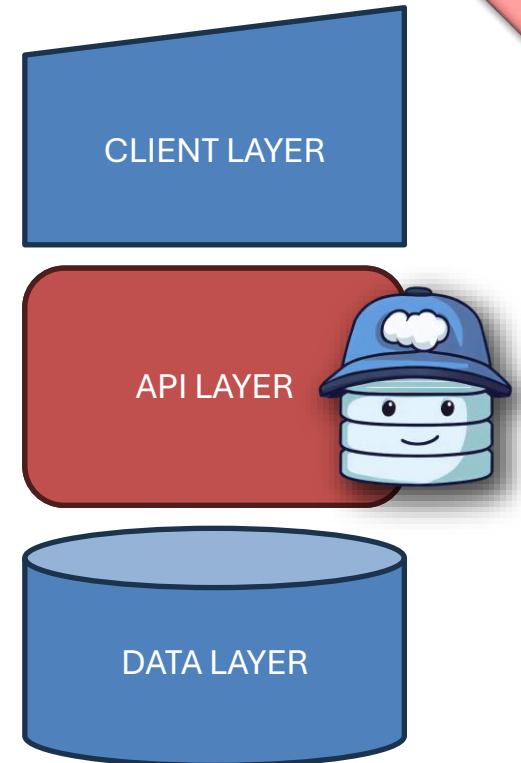
Data API builder

- Open Source and Free (<https://aka.ms/dab>)
- Container-based
- Zero code
 - Configuration-driven
- Supported authentication:
 - EasyAuth, EntralID/AAD, JWT tokens
- Supported endpoints:
 - Standard REST API and Standard GraphQL
- Supported databases:
 - Azure SQL DB & SQL Server
 - Cosmos DB (document)
 - PostgreSQL
 - MySQL



Data API builder

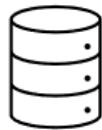
- Built in developer tools and clients
 - OpenAPI, Swagger, Banana Cake Pop
 - DAB CLI
- Database Objects Supported
 - Tables
 - Views
 - Stored Procedures
- GraphQL Relationships
 - Foreign Keys
 - One-to-Many/Many-to-One/Many-to-Many
- Authorization
 - Database Context
 - Role-based authorization using received claims
 - Item-level security via policy expressions



"Plain"
GraphQL
Schema

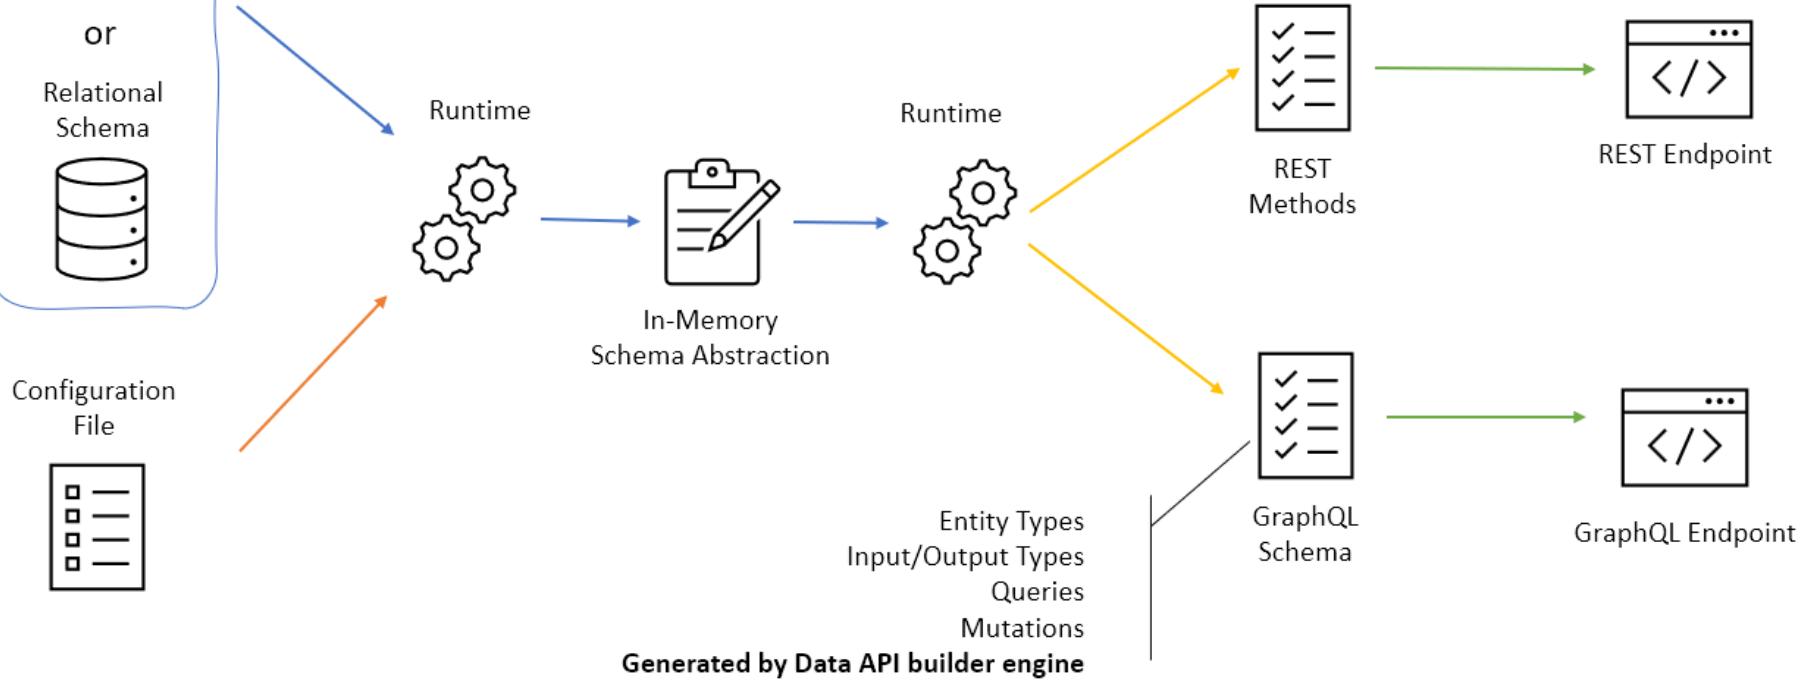


or
Relational
Schema



Entity Types only
Provided by Developer

Data API builder



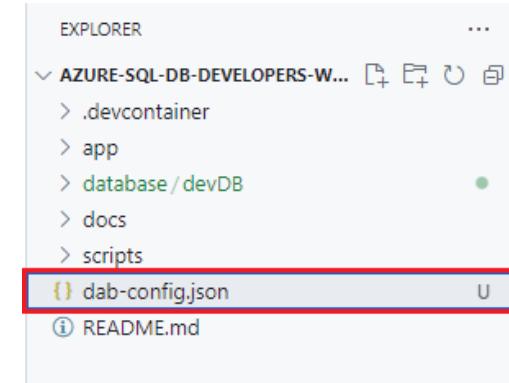
Chapter 3: Data API Builder

- Setup Data API Builder
- Create REST/GraphQL entities
- Test the endpoints

Setup Data API Builder

Create the Data API Builder initialization file

```
dab init --database-type "mssql" --connection-string  
"Server=localhost;Database=devDB;User ID=vscode;Password=  
'XXXXXX';TrustServerCertificate=true" --host-mode  
"Development" --rest.path "rest" --set-session-context  
true
```

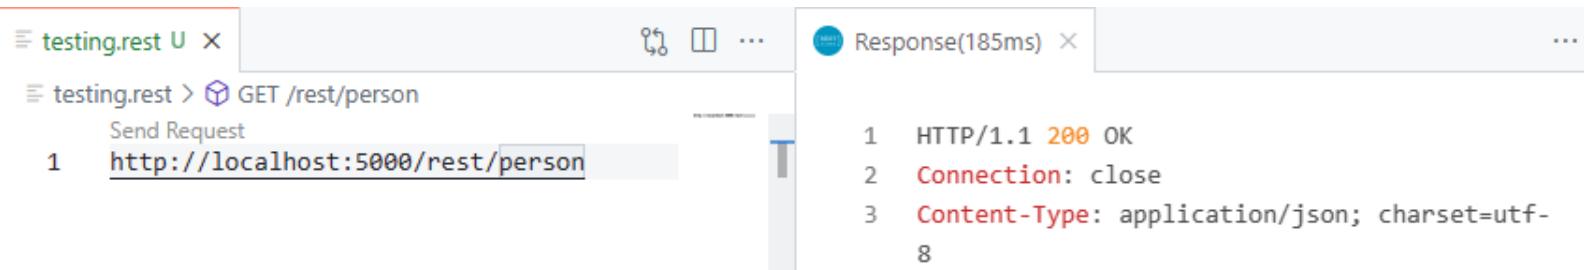


Create REST/GraphQL entities

Use DAB CLI to add the tables, stored procedures, and relationships to the dab-config.json file

- **dab add person --source dbo.person --permissions "anonymous:***"
- **dab update person --relationship "address" --cardinality "many" --target.entity "address"**
- **dab add getPersonByPet --source dbo.get_person_by_pet --source.type "stored-procedure" --source.params "pet:" --permissions "anonymous:execute" --rest.methods "get" --graphql.operation "query"**

Test the endpoints



The screenshot shows a browser window with two tabs. The left tab is titled 'testing.rest' and contains a request for 'GET /rest/person' with the URL 'http://localhost:5000/rest/person'. The right tab is titled 'Response(185ms)' and displays the JSON response. The response includes standard HTTP headers and a JSON payload containing a single person object.

```
1 HTTP/1.1 200 OK
2 Connection: close
3 Content-Type: application/json; charset=utf-8
4 Date: Thu, 19 Oct 2023 21:06:45 GMT
5 Server: Kestrel
6 Transfer-Encoding: chunked
7 x-ms-correlation-id: 28f732dc-081d-41b9-b809-ee655d33c3b5
8
9 {
10   "value": [
11     {
12       "person_id": 1,
13       "person_name": "Bill",
14       "person_email": "bill@contoso.com",
15       "pet_preference": "Dogs"
16     },
17   ]
}
```

Use codespaces to test your REST/GraphQL endpoints

Knowledge Check

Data API builder lets you REST/GraphQL enable?

- A. Tables
- B. Views
- C. Stored Procedures
- D. A and B
- E. A, B, and C

On average, how many lines of code do I need to write to GraphQL enable a database object with Data API Builder

- A. 77
- B. 4000
- C. None
- D. 4 or 5

Chapter 3: Data API Builder

- Setup Data API Builder
- Create REST/GraphQL entities
- Test the endpoints

A photograph of a large, brown and white adult llama standing on the left, facing right. A smaller, white cria stands next to it, looking up at the adult. They are in a grassy field with other llamas in the background under a clear sky.

START CHAPTER 3

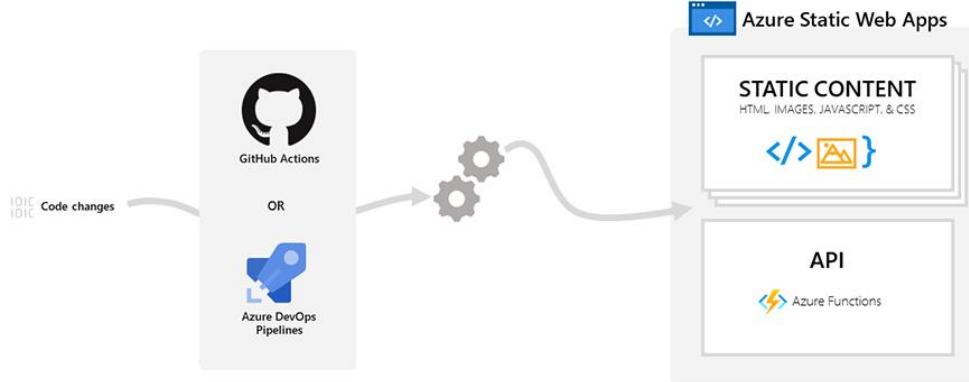
CHAPTER 4: AZURE STATIC WEB APPS



Azure Static Web Apps

Automatically build and deploy full stack web apps to Azure from a code repository.

- **Build modern web applications** with JavaScript frameworks and libraries like [Angular](#), [React](#), [Svelte](#), [Vue](#), or using [Blazor](#) to create WebAssembly applications, with an [Azure Functions](#) back-end.
- **Publish static sites** with frameworks like [Gatsby](#), [Hugo](#), [VuePress](#).
- **Deploy web applications** with frameworks like [Next.js](#) and [Nuxt.js](#).



Key Features

- **Web hosting** for static content like HTML, CSS, JavaScript, and images.
- **Integrated API** support provided by managed Azure Functions, with the option to link an existing function app, web app, container app, or API Management instance using a standard account. If you need your API in a region that doesn't support [managed functions](#), you can [bring your own functions](#) to your app.
- **First-class GitHub and Azure DevOps integration** that allows repository changes to trigger builds and deployments.
- **Globally distributed** static content, putting content closer to your users.
- **Free SSL certificates**, which are automatically renewed.
- **Custom domains** to provide branded customizations to your app.

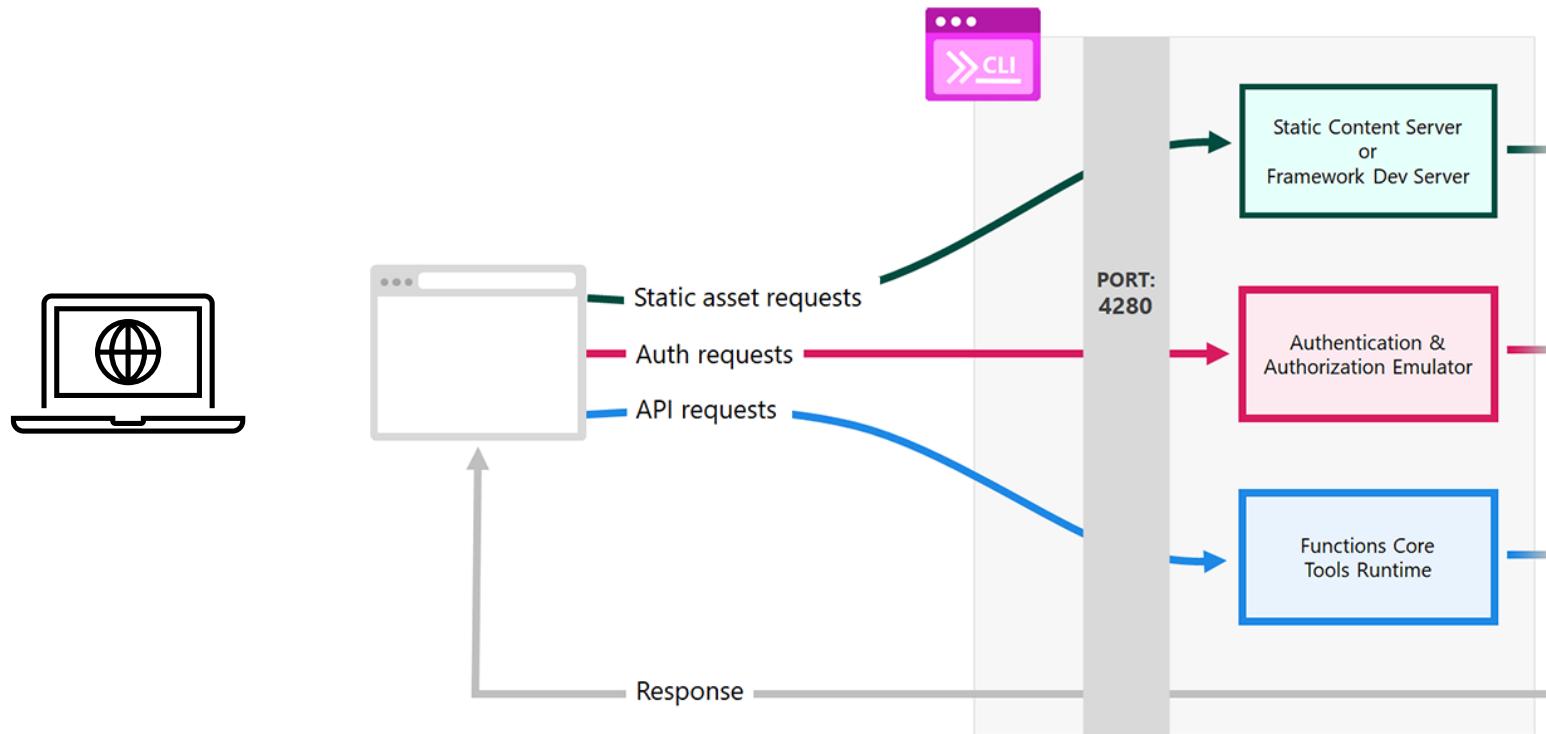
Key Features

- **Seamless security model** with a reverse-proxy when calling APIs, which requires no CORS configuration.
- **Authentication provider integrations** with Microsoft Entra ID and GitHub.
- **Customizable authorization role definition** and assignments.
- **Back-end routing rules** enabling full control over the content and routes you serve.
- **Generated staging versions** powered by pull requests enabling preview versions of your site before publishing.
- **CLI support** through the [Azure CLI](#) to create cloud resources, and via the [Azure Static Web Apps CLI](#) for local development.

Azure Static Web Apps CLI

- Serve static app assets, or proxy to your app dev server
- Serve API requests, or proxy to APIs running in Azure Functions Core Tools
- Emulate authentication and authorization
- Emulate Static Web Apps configuration, including routing and ACL roles
- Deploy your app to Azure Static Web Apps

Azure Static Web Apps CLI



Chapter 4: SWA

- Initialize SWA
- Test the web applications
- Use SWA authentication with Data API builder

But First....

Creating a mapping in Data API Builder

```
dab update todo --map "position:order"
```

To map a UI element called order to the position column (in dab-config.json)

```
"mappings": {  
    "position": "order"  
}
```

Initialize SWA and Test

Create the connections file and start with the sample JavaScript application

- 1) cp dab-config.json ./swa-db-connections/staticwebapp.database.config.json
- 2) swa start --app-location ./app --data-api-location ./swa-db-connections

Sample JavaScript App

Pet Preference:

Submit

Person ID	Person Name	Person Email	Pet Preference
1	Bill	bill@contoso.com	Dogs
2	Frank	frank@contoso.com	Cats
3	Riley	Riley@contoso.com	Cats

[Login](#)

Use SWA authentication with Data API builder

In the permissions section, change the role from **anonymous** to **authenticated**

```
"permissions": [  
    {  
        "role": "authenticated",  
        "actions": [  
            "execute"  
        ]  
    }  
],
```

Username

Fluffy Bunny



Username or email address of the user

User's roles

anonymous

authenticated

Roles used during authorization. One role per line.

Note: roles "authenticated" and "anonymous" will be added automatically if not provided.

User's claims

[]

Claims from the identity provider. JSON array of claims. See [documentation](#) for example claims.

Login

Clear

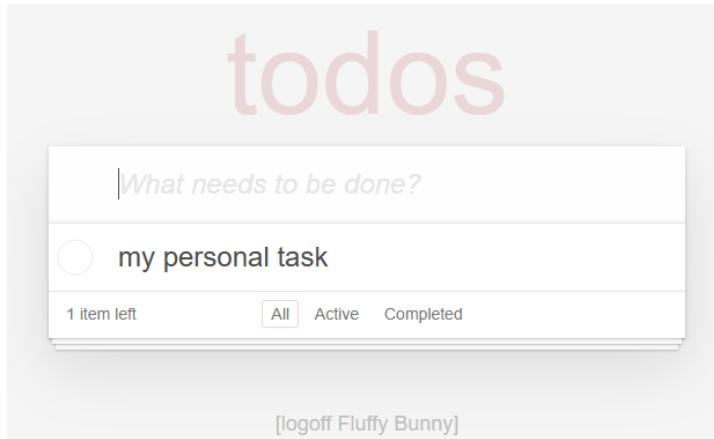
Working with the ToDo Application

- Change the `swa-cli.config.json`
- Build the app (`swa build`)
- Start swa and test the ToDo application

```
swa start --data-api-location ./swa-db-connections
```

Working with the ToDo Application

In the permissions section, change replace the permissions section for the Todo table with the following code:



```
"permissions": [
  {
    "role": "anonymous",
    "actions": [
      {
        "action": "*",
        "policy": {
          "database": "@item.owner_id eq 'public'"
        }
      }
    ]
  },
  {
    "role": "authenticated",
    "actions": [
      {
        "action": "*",
        "policy": {
          "database": "@item.owner_id eq @claims.userId"
        }
      }
    ]
  }
],
```

Working with the ToDo Application

Move from the REST endpoints to stored procedures

1. mv client/src/components/ToDoList.vue
labFiles/ToDoList.vue.DAB1
2. mv labFiles/ToDoList.vue.SP
client/src/components/ToDoList.vue
3. swa start --data-api-location ./swa-db-connections

Knowledge Check

Static Web Apps is a cloud only service/tool?

A. True

B. False

A major issue with using Static Web Apps is that I have to deploy an instance in each region I want to use?

A. True

B. False

Chapter 4: SWA

- Initialize SWA
- Test the web applications
- Use SWA authentication with Data API builder



START CHAPTER 4

Announcing...

Azure Database Fleet Manager - Private Preview

- New service designed to help SaaS solution builders create multi-tenant database tiers.
 - Offers a simple and intuitive API to provision and manage tenant databases, seamlessly scaling from few to millions.
 - Tenants are organized in tiers. A tier represents a grouping of settings the platform enforces on all databases in a tier.
 - Fleet Manager dynamically load balances resource utilization for databases in a tier, optimizing costs and performance for all tenants.
- Available with Azure SQL Database Elastic Pools and single databases now, more options will follow.
- More details: <https://aka.ms/dbfleetmanager-prpr>

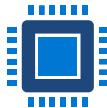
CHAPTER 5:

DEPLOY TO AZURE SQL DATABASE



Azure SQL Database free of charge

Get 100,000 vCore seconds of serverless compute and 32 GB of storage every month!



Azure SQL Database with serverless compute

Flexible compute automatically scales to meet demand.

No time limits

Apply this free offer for the life of your subscription.

Need more? No problem.

Stick with the default auto-pause option or continue usage for additional charges.

What's included:



One Azure SQL Database with serverless compute per Azure Subscription with 100,000 vCore seconds every month.



32 GB data storage +
32 GB backup storage.

Learn More: aka.ms/sqlfreeoffer



Want to try Azure SQL Database for free? Create a free serverless database with the first 100,000 vCore seconds, 32GB of data, and 32GB of backup storage free per month for the lifetime of the subscription. [Learn more](#)

Apply offer (Preview)

Extra Credit

AZURE DATABASE MIGRATION SERVICE

Azure Database Migration Service

Migrate your databases to Azure at scale



A fully managed service designed to streamline the migration of SQL Server databases, whether running on-premises or in other cloud environments



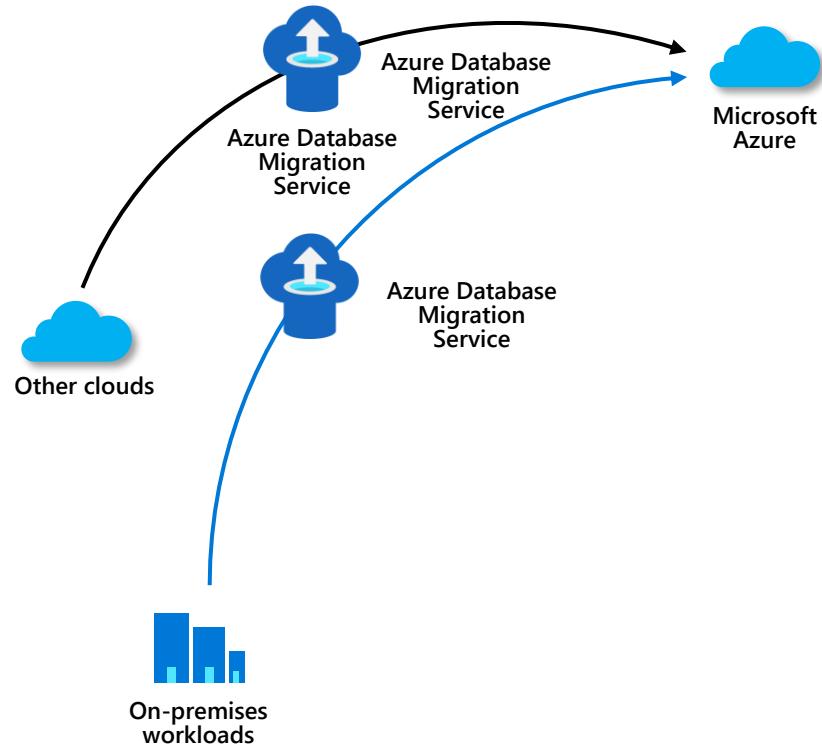
Highly resilient and self-healing migration experience, with near-zero downtime



Seamlessly migrate to Azure using familiar tools like the Azure Portal and Azure Data Studio.



Automate with ease using PowerShell and Azure CLI



On-premises

Microsoft®
SQL Server®



mongoDB®



Google Cloud

Microsoft Azure



Azure Database Migration Service

Assessment



Migration

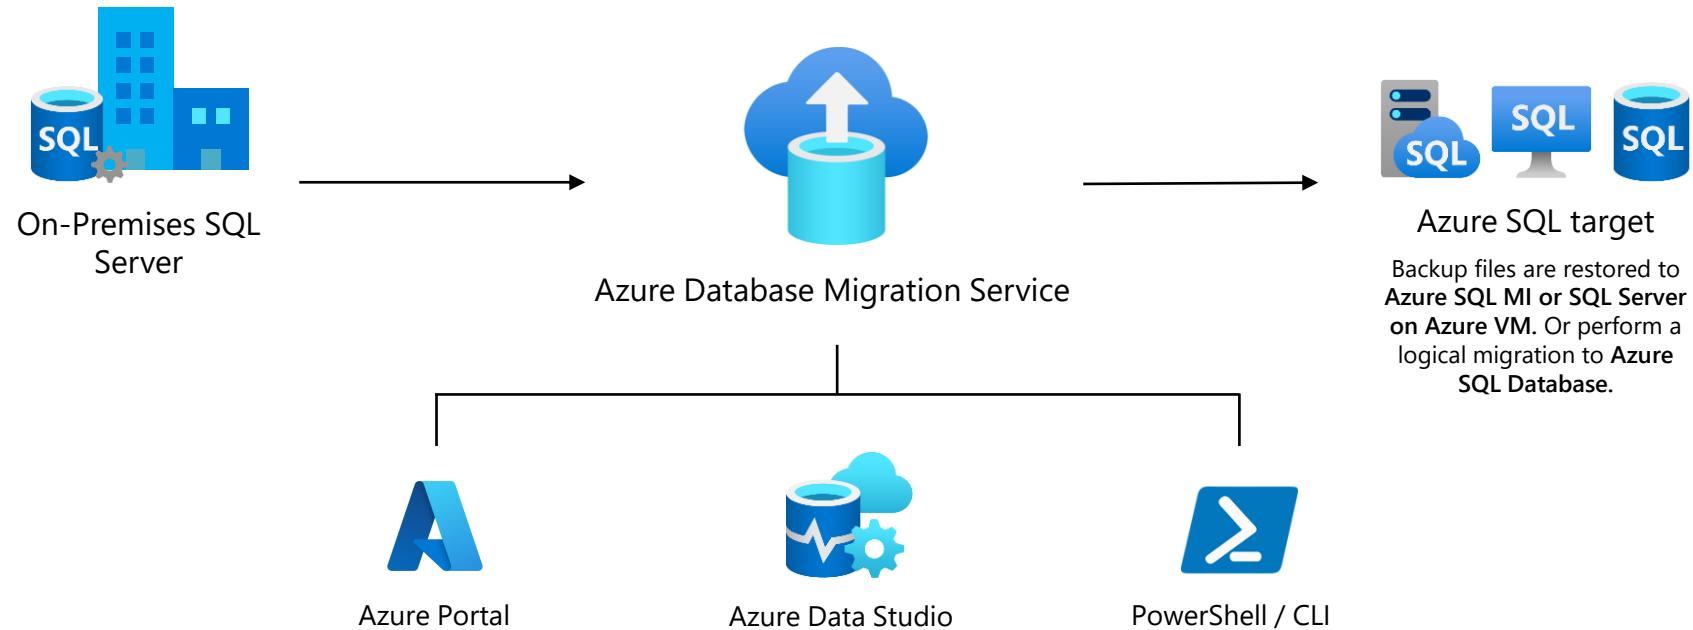


Azure Database
Migration Service

Azure DMS
(ADS, Azure portal, PowerShell, CLI)
Azure PostgreSQL Migration (ADS)
Azure Migrate

Azure Database Migration Service

Migration tools



Chapter 5: Deploy to Azure

- Create a Free Azure SQL Database
- Connect to the Azure SQL Database
- Publish objects to the Azure SQL Database

Create a Free Azure SQL Database

- Login to the Azure Portal
- Create a new SQL Database
- Apply the promotion
- Create!

The screenshot shows the Azure portal interface for creating a new SQL database. At the top, there's a large blue 'SQL' icon. Below it, the 'Cost summary' section displays the following details:

General Purpose (GP_S_Gen5_2)	
Cost per GB (in USD)	0.00
Max storage selected (in GB)	x 41.6
First 32 GB storage free	
First 100,000 vCore seconds free	
Overage billing ¹	Disabled
ESTIMATED STORAGE COST / MONTH	0.00 usd
COMPUTE COST / VCORE SECOND ²	0.000000 usd

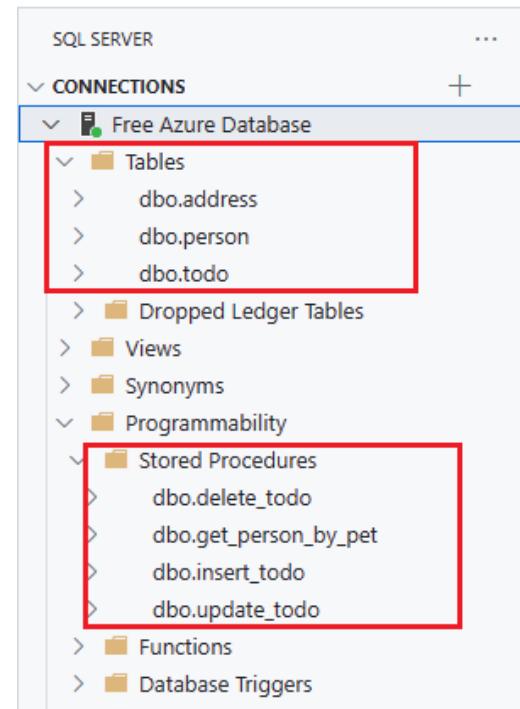
At the bottom, two footnotes provide additional information:

¹ There will be no charges for usage within the free limits. The database will be paused automatically when the free limits are reached.

² Serverless databases are billed in vCore seconds based on a combination of CPU and memory utilization. [Learn more about serverless billing](#)

Connect and Publish to the Azure SQL Database

- Create a connection in codespace
- Deploy the database project to the free Azure SQL Database
- Verify the deployment and post-install script



Knowledge Check

Azure SQL Database is expensive as a development platform?

- A. True
- B. False

Chapter 5: Deploy to Azure

- Create a Free Azure SQL Database
- Connect to the Azure SQL Database
- Publish objects to the Azure SQL Database

START CHAPTER 5



CHAPTER 6: INVOKE REST ENDPOINTS



Services Integration

I really want to use other Azure products and services from within my Azure SQL Database without having to write complex code, functions, or scheduled events. Is this even possible?

External REST Endpoint Invocation

Using a single stored procedure, call REST/GraphQL service endpoints in Azure directly from within an Azure SQL Database.



Use Cases

- Activate workflows
- Data enrichment
- Cache invalidation / update
- Augment business/process logic
- Update websites
- Integrate with event-based architectures
- Data streams

Security Included

In the Database

- Requires EXECUTE ANY EXTERNAL ENDPOINT database permission.

Endpoint Authentication

- Header
- Query String/URL
- Managed Identity

Easy to use

One simple and easy call to a REST endpoint directly from the Database

```
EXEC @returnValue = sp_invoke_external_rest_endpoint
[ @url = ] N'url'
[ , [ @payload = ] N'request_payload' ]
[ , [ @headers = ] N'http_headers_as_json_array' ]
[ , [ @method = ] 'GET' | 'POST' | 'PUT' | 'PATCH' | 'DELETE' | 'HEAD' ]
[ , [ @timeout = ] seconds ]
[ , [ @credential = ] credential ]
[ , @response OUTPUT ]
```

Chapter 6: Invoke REST endpoints

- Invoke an Azure Function endpoint
- Invoke an OpenAI endpoint
- Add translation logic to the todo application with OpenAI

Invoke an Azure Function endpoint

- Use a function to augment/enhance data
- Get the conversion rate for currency X and see what the price is
- Call endpoint and extract rate from JSON into a sql variable

```
set @priceConversion = (select JSON_VALUE(@response,'$.result.priceConversion') AS priceConversion);
```

```
select product_id, product_name, ListPrice,  
cast(round(ListPrice*@priceConversion,2,1) as money) AS  
convertedPriceInYen from dbo.products
```

Invoke an OpenAI endpoint

Use AI to get generated content from data in the database

```
set @adcopy =
(select 'You are an experienced marketing expert named Don Chase
Katz. Generate 200 letters of ad copy to '
+ person_name +
' to convince them to love ' +
case when pet_preference = 'Cats' then 'Dogs'
      when pet_preference = 'Dogs' then 'Cats'
end
from person
where person_id = 1);
```

Point SWA at the Azure SQL Database

Change the connection string in the
staticwebapp.database.config.json file to
point to the Azure SQL Database



```
{ staticwebapp.database.config.json 3, U X
swa-db-connections > staticwebapp.database.config.json > data-source
1 {
2   "$schema": "https://github.com/Azure/data-api-builder/releases/download/v0.8.52/dab.draft.schema.json",
3   "data-source": {
4     "database-type": "mssql",
5     "connection-string": "Server=freedbsqlserver.database.windows.net;Database=freeDB;User ID=sqladmin;Password=PASSWORD;TrustServerCertificate=true",
6     "options": {
7       "set-session-context": false
8     }
9   },

```

Add translations with OpenAI

Alter the `insert_todo` stored procedure to call OpenAI to translate the text of a ToDo before an insert is done

```
declare @payload nvarchar(max) =  
N'{"messages": [{"role": "system", "content": "Translate  
\""+(@title)+"\" into german, only respond with the  
translation"}]}'
```

Knowledge Check

When calling REST endpoints from the database, the authentication method(s) not available is?

- A. Managed Identity
- B. Mangled Identity
- C. Header
- D. Query String/URL
- E. B and D

Chapter 6: Invoke REST endpoints

- Invoke an Azure Function endpoint
- Invoke an OpenAI endpoint
- Add translation logic to the todo application with OpenAI

A snowy owl is perched on a light-colored, textured surface, possibly snow or sand. The owl's white face and chest are prominent, with its eyes closed and its beak wide open, revealing a bright red gape. Its dark, barred feathers are visible on its back and wings. The background is a uniform, light gray.

START CHAPTER 6

CHAPTER 7: AZURE SQL BINDINGS



Sup?

Azure Functions Core Tools

- Develop Azure Functions locally using Core Tools
- Deploy your code project to Azure
 - Function, App Container,
- Work with app settings locally
(local.settings.json)

Change Detection & Change Stream

How can I leverage the change tracking features in Azure SQL Database (or SQL Server/MI/VM) to expose data changes via an outgoing change stream without the need for custom code and scheduler functions?

Azure SQL bindings for Azure Functions

- **Azure SQL**
 - encompasses all Azure SQL products (DB, MI, VM)
 - + SQL Server with cloud connectivity
- **Azure Functions**
 - serverless runtime for standalone use or integrated with Azure Static Web Apps
- **Popular Languages**
 - C# (in proc and out of proc), JavaScript, Python, Java, and PowerShell

Three Azure SQL bindings for Azure Functions

Input binding



object in function



SQL query results



Azure SQL Input bindings take a SQL query or stored procedure to run and returns the output to the function.

Output binding



object in function



SQL table



Azure SQL Output bindings take a list of rows and upserts them to the user table.

SQL trigger



SQL data change



starts function



The Azure SQL trigger uses SQL change tracking functionality to monitor a SQL table for changes

Azure SQL bindings for Azure Functions

Use Cases

- **Data enrichment**

- Determine if a value is an outlier or not using Azure Cognitive services/OpenAI
- Perform reverse geocoding using Azure Functions
- Call a REST/GraphQL service from a function with data from Azure SQL

- **Start complex processing**

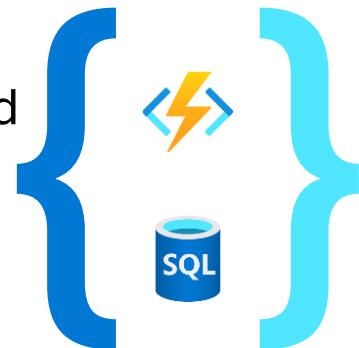
- Call a Function to kick off some complex process and return the data



Azure SQL bindings for Azure Functions

Use Cases

- **Update websites**
 - Broadcast a SignalR message
- **Integrate with event-based architectures**
 - Send data to Event Hubs for further integration options
- **Create a change data stream**
 - Send data to Stream Analytics for further investigation/fraud detection



SQL bindings + Azure SQL compatibility

- The bindings use the [OPENJSON](#) statement which requires a **database compatibility level of 130 or higher (2016 or higher)**.
- Databases on **SQL Server, Azure SQL Database, or Azure SQL Managed Instance** which meet the compatibility level requirement above are supported.

Connecting to the database

- SQL bindings connect to the target database by using a **Connection String configured in the app settings**. This will require a login be created that the function will use to access the server.
- For local testing and development using a SQL (username/password) or Azure Active Directory Login is typically the easiest, but for deployed function apps it is **recommended to use Azure Active Directory Managed Authentication**.

SQL trigger app settings

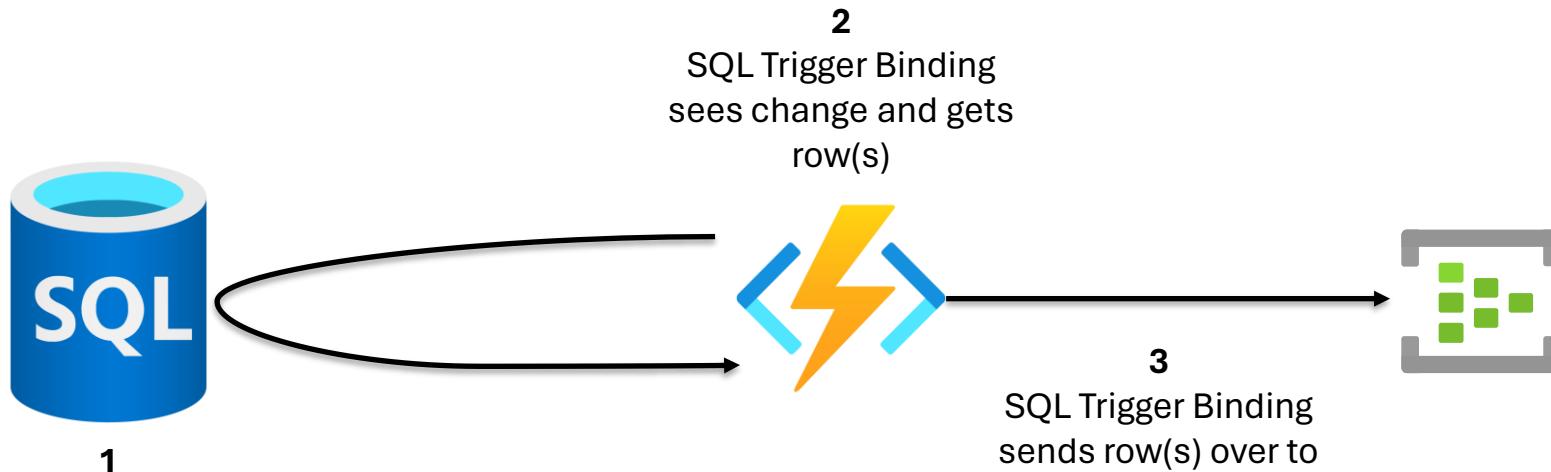
- **Sql_Trigger_BatchSize**: controls the number of changes processed at once before being sent to the triggered function.
- **Sql_Trigger_PollingIntervalMs**: controls the delay in milliseconds between processing each batch of changes.
- **Sql_Trigger_MaxChangesPerWorker**: controls the upper limit on the number of pending changes in the user table that are allowed per application-worker. If the count of changes exceeds this limit, it may result in a scale out. (only applies for Azure Function Apps with runtime driven scaling enabled)

Change stream scenario

Requirement is to capture changes from a table and put them onto an Azure Event Hub for down stream subscribers (microservices).

Also, I don't want to write a lot of code.

Change stream architecture



```
1  
Update on Products table  
  UPDATE products  
  SET product_price = 3.99  
  WHERE product_id = 2;
```

Chapter 7: Azure SQL Bindings

- Create a function project
- Add a SQL trigger binding
- Run the function and test on the Azure SQL Database

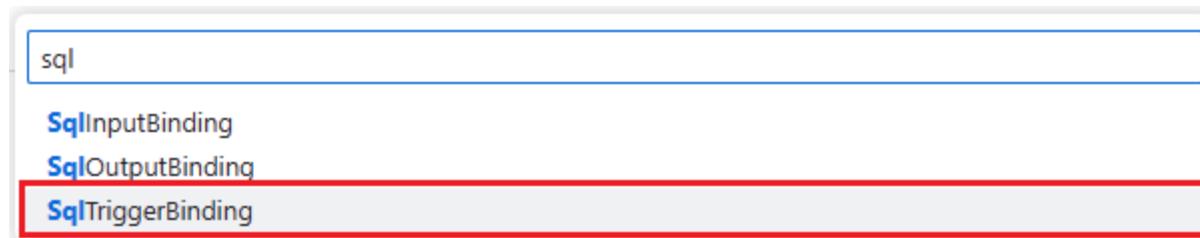
Create a function project

- Create a .NET function using core tools

```
func init triggerBinding --worker-runtime dotnet
```

Add a SQL trigger binding

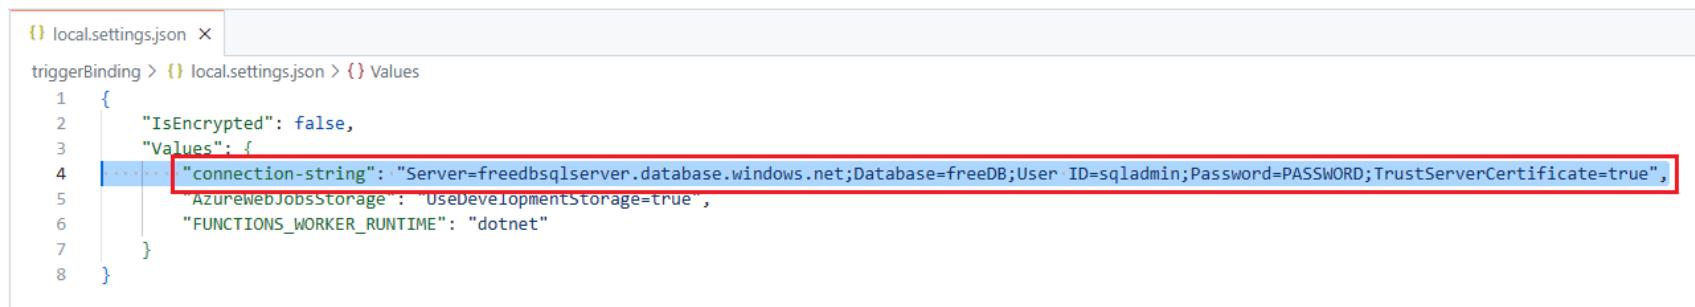
- Use the Create Function wizard to create the SQL Trigger Binding
 - can also be used to create input and output bindings



Run and Test the Function

- Use core tools to run locally
- Connect to a local or cloud SQL Database

```
func host start
```



```
local.settings.json x
triggerBinding > local.settings.json > Values
1 {
2     "IsEncrypted": false,
3     "Values": [
4         {
5             "connection-string": "Server=freedbsqlserver.database.windows.net;Database=freeDB;User ID$sqladmin;Password=PASSWORD;TrustServerCertificate=true",
6                 "AzureWebJobsStorage": "UseDevelopmentStorage=true",
7                 "FUNCTIONS_WORKER_RUNTIME": "dotnet"
8         }

```

Knowledge Check

Functions can only be developed in the cloud?

A. True

B. False

Azure SQL bindings for Azure Functions can only be used with Azure SQL Database?

A. True

B. False

Chapter 7: Azure SQL Bindings

- Create a function project
- Add a SQL trigger binding
- Run the function and test on the Azure SQL Database

A large flock of Barn Swallows is perched on several horizontal wires against a clear blue sky. The birds are arranged in four distinct rows. The top row has 15 birds, the second row has 18, the third row has 12, and the bottom row has 10. Most birds are facing right, while a few are facing left or have their heads turned. The wires are thin and light-colored, stretching across the frame.

START CHAPTER 7

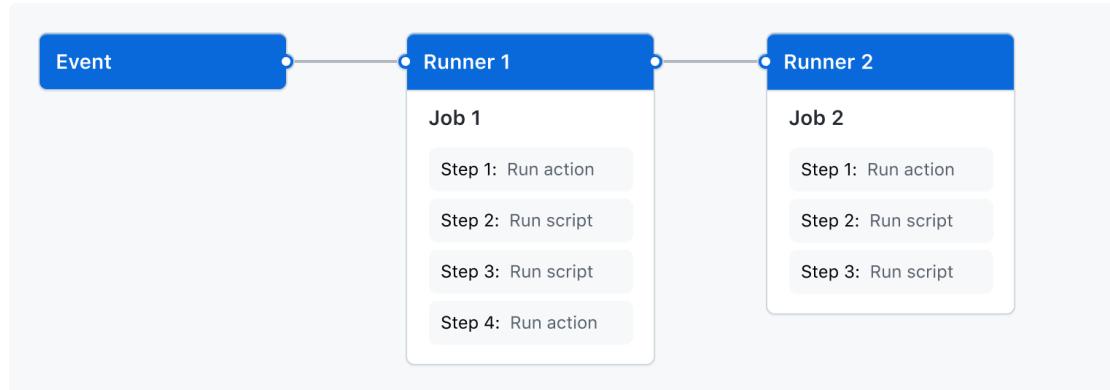
CHAPTER 8:

CI/CD WITH GITHUB ACTIONS



GitHub Actions

- Are a continuous integration and continuous delivery (CI/CD) platform
- Allows you to run workflows based on repository events
- Test or deploy code with workflows

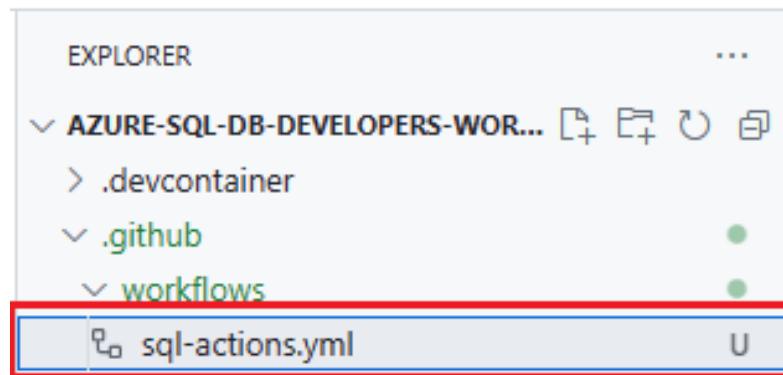


Chapter 8: CI/CD with GitHub Actions

- Create a workflow folder and actions file
- Use sql-actions for DB creation and testing
- Push changes to start the action

Create a workflow folder

- GitHub actions are in the .github/workflows folder
- Can contain multiple action workflows



Use sql-actions

- Sql-actions are SQL specific activities that can be run in a workflow
- Publish (.dacpac) and run scripts on database
- Integrated with GitHub Secrets

Push to Start

- Push code back to your fork to start the workflow
- Workflow will:
 - Create a SQL 2022 DB in docker
 - Create a test DB and set server parameters
 - Deploy the database project to the test database
 - Log all steps for review
 - Automatically deletes environment when done

sql-actions.yml

```
# .github/workflows/sql-actions.yml  
  
name: SQL Server container in deployment  
test pipeline  
on: [push]
```

Name of the workflow

Runs on a git push

sql-actions.yml

```
jobs:  
  build-and-deploy:  
    # Containers must run in Linux based operating  
    systems  
    runs-on: ubuntu-latest  
  
    # service/sidecar container for azure-sql-2022  
    services:  
      mssql:  
        image: mcr.microsoft.com/mssql/server:2022-  
latest  
        env:  
          ACCEPT_EULA: 1  
          SA_PASSWORD: XXXXXXXX  
        ports:  
          - 1433:1433
```

Groups together all the jobs that run in the workflow.

OS the container is going to run in

A docker container that will run SQL 2022

sql-actions.yml

```
steps:
  - name: 'Checkout GitHub Action'
    uses: actions/checkout@v4

  - name: 'wait for sql to be ready'
    run: |
      set +o pipefail +e
      for i in {1..60};
      do
        sqlcmd -S localhost -U sa -P XXXXXXXX -d master -Q "running
"select getdate()"
        if [ $? -eq 0 ]
        then
          echo "sql server ready"
          break
        else
          echo "not ready yet..."
          sleep 1
        fi
      done
      set -o pipefail -e
```

Check out the repository code

Wait for the database to be up and running

sql-actions.yml

```
- name: 'Create and setup database'
  uses: azure/sql-action@v2
  with:
    connection-string: "Server=localhost;Initial
Catalog=master;User ID=sa;Password=
XXXXXXXX;Encrypt=False;TrustServerCertificate=False;" # the local
connection string
    path: './labFiles/setupDatabase.sql' # the sql script to
create db and configure for clr

- name: 'Deploy Database Project'
  uses: azure/sql-action@v2.2
  with:
    connection-string: "Server=localhost;Initial
Catalog=testingDB;User ID=sa;Password=
XXXXXXXX;Encrypt=False;TrustServerCertificate=False;" # the local
connection string
    path: './database/devDB/devDB.sqlproj' # the SQLproj file
    action: 'Publish'
```

Once the DB is ready, run the following script.
(creates testDB and run env config options)

Deploy the database project we created to the testDB

Knowledge Check

If I use a developer edition of SQL Server 2022 in my GitHub action, it will cost me?

- A. \$100,000,000
- B. \$15
- C. \$0
- D. The cost of a cup of coffee a day for 30 days

Chapter 8: CI/CD with GitHub Actions

- Create a workflow folder and actions file
- Use sql-actions for DB creation and testing
- Push changes to start the action

START CHAPTER 8



THANK YOU!



Session evaluation

Your feedback is important to us



Evaluate this session at:

www.PASSDataCommunitySummit.com/evaluation

